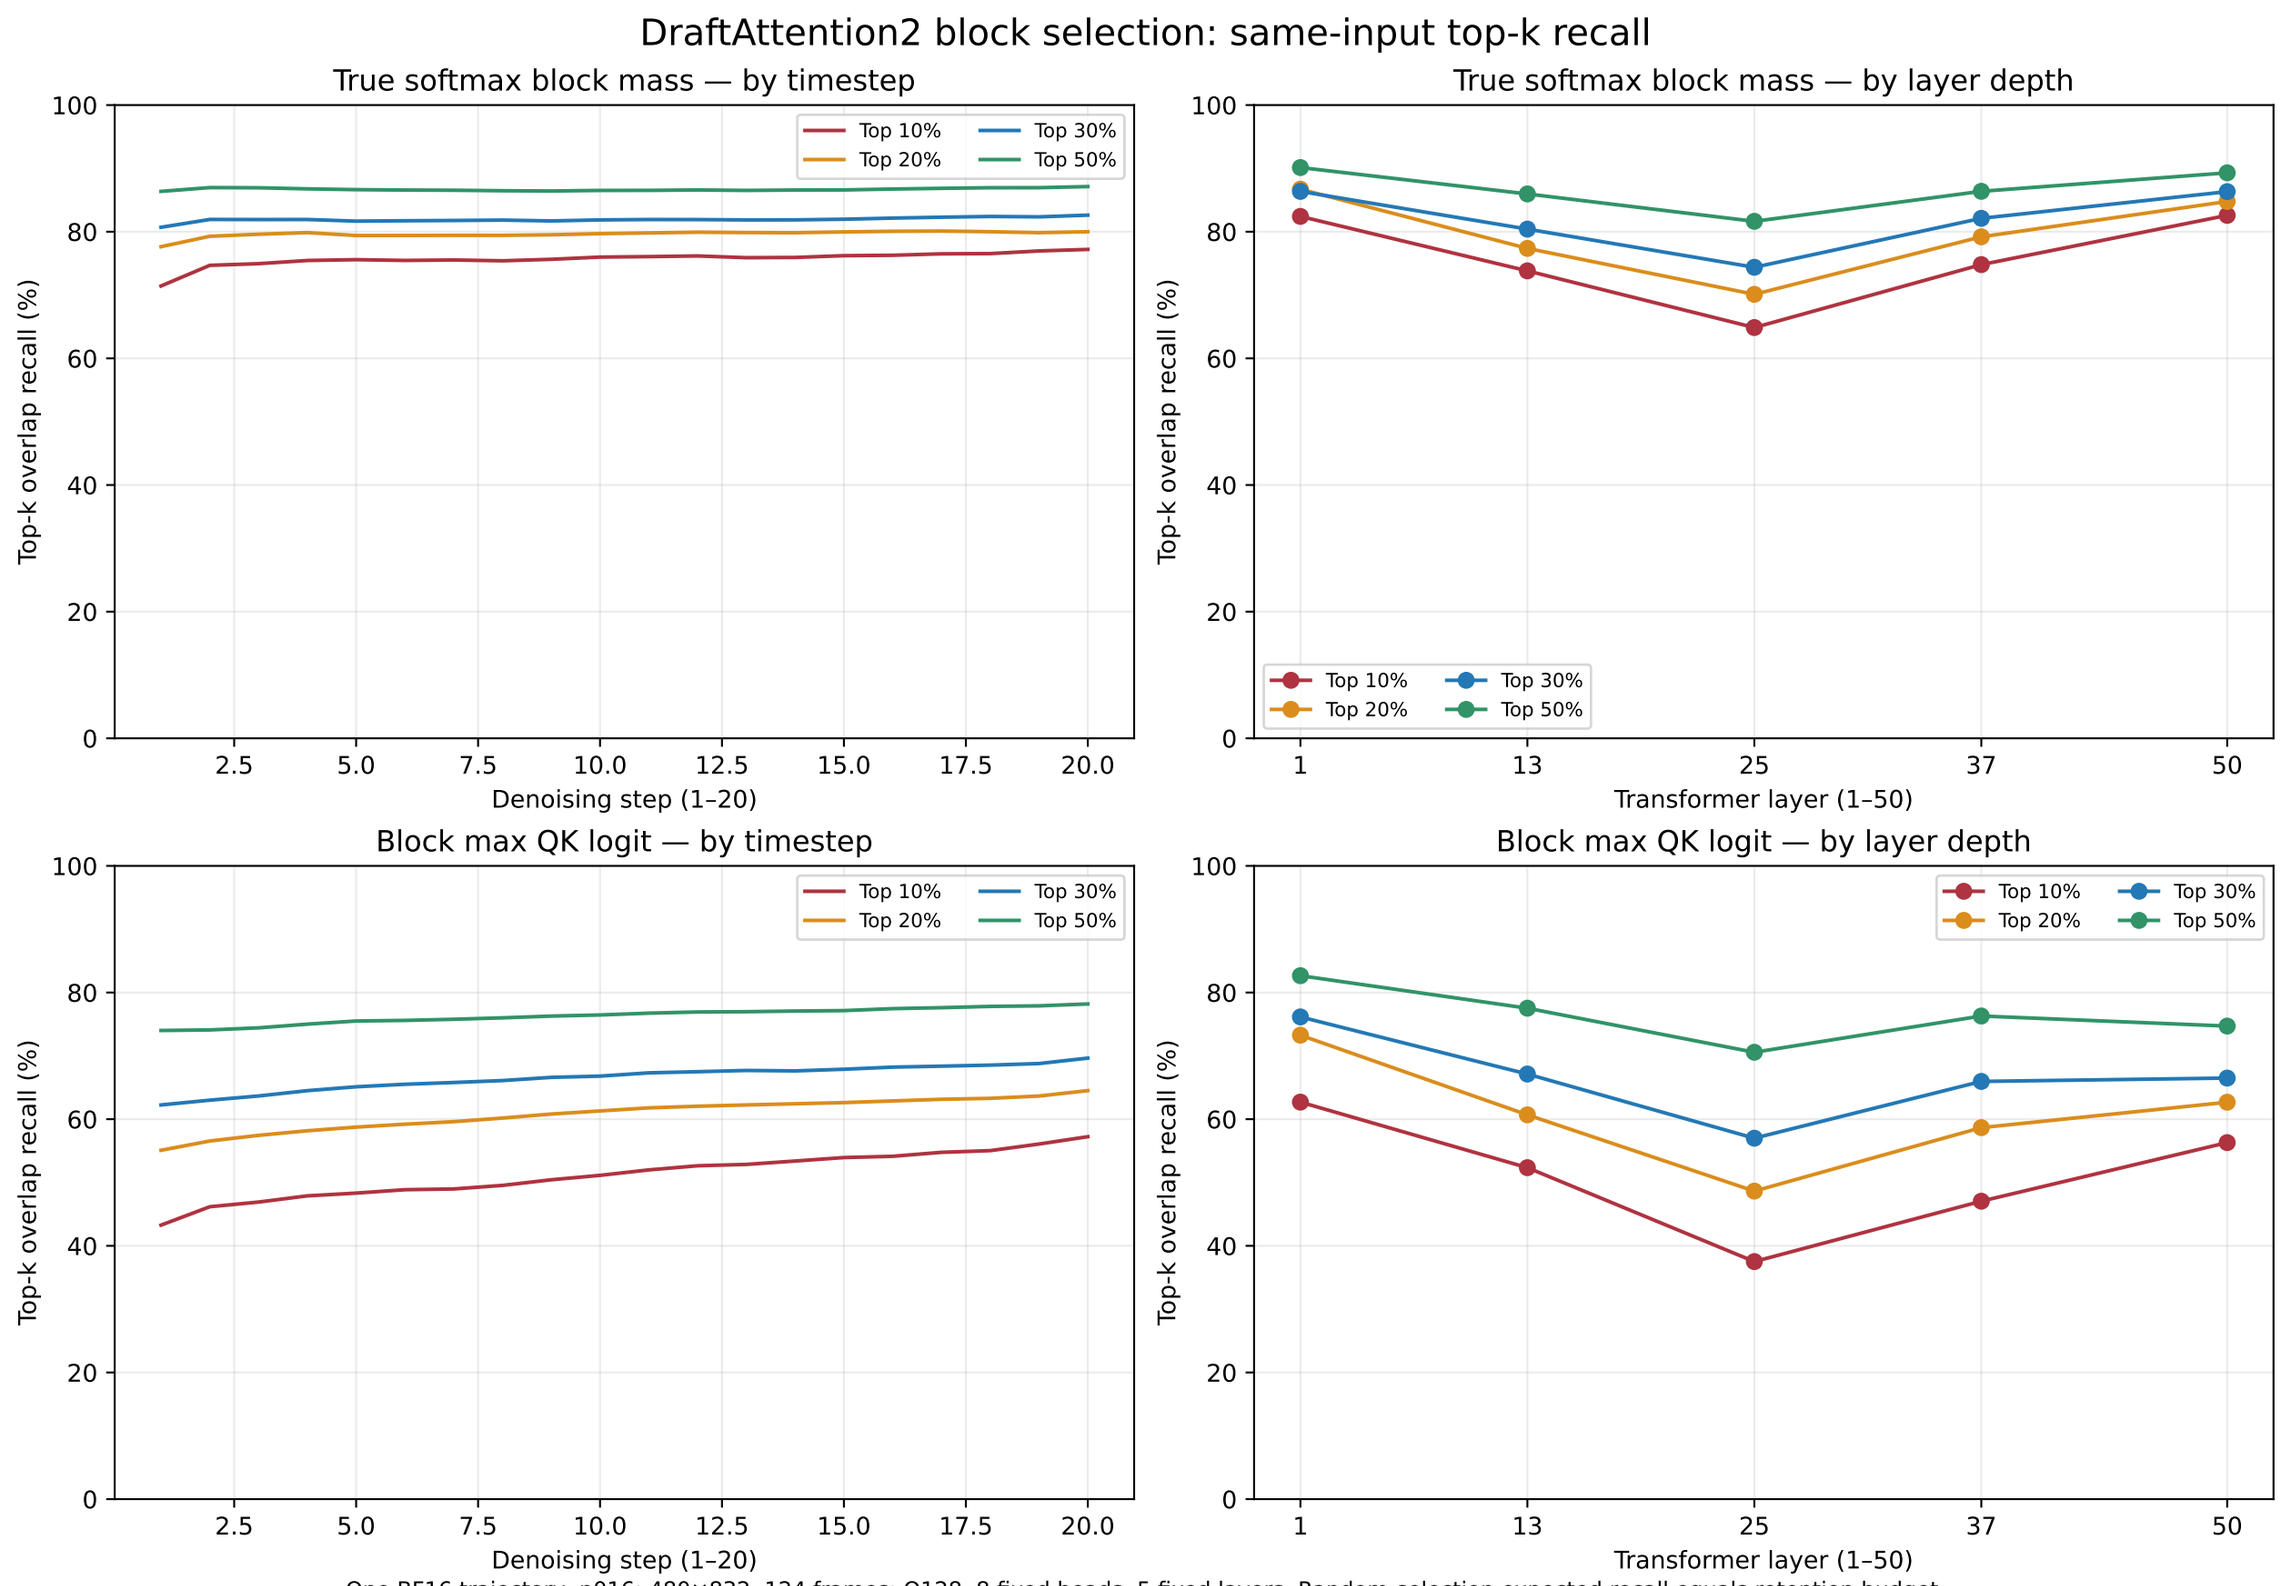
DraftAttention2 block selection: same-input top-k recall
True softmax block mass — by timestep
Top-k overlap recall (%)
Denoising step (1–20)
Top 10%
Top 20%
Top 30%
Top 50%
True softmax block mass — by layer depth
Transformer layer (1–50)
Block max QK logit — by timestep
Block max QK logit — by layer depth
One BF16 trajectory, p016; 480×832, 124 frames; Q128, 8 fixed heads, 5 fixed layers. Random-selection expected recall equals retention budget.