DraftAttention2 block selection: same-input top-k recall

True softmax block mass — by timestep

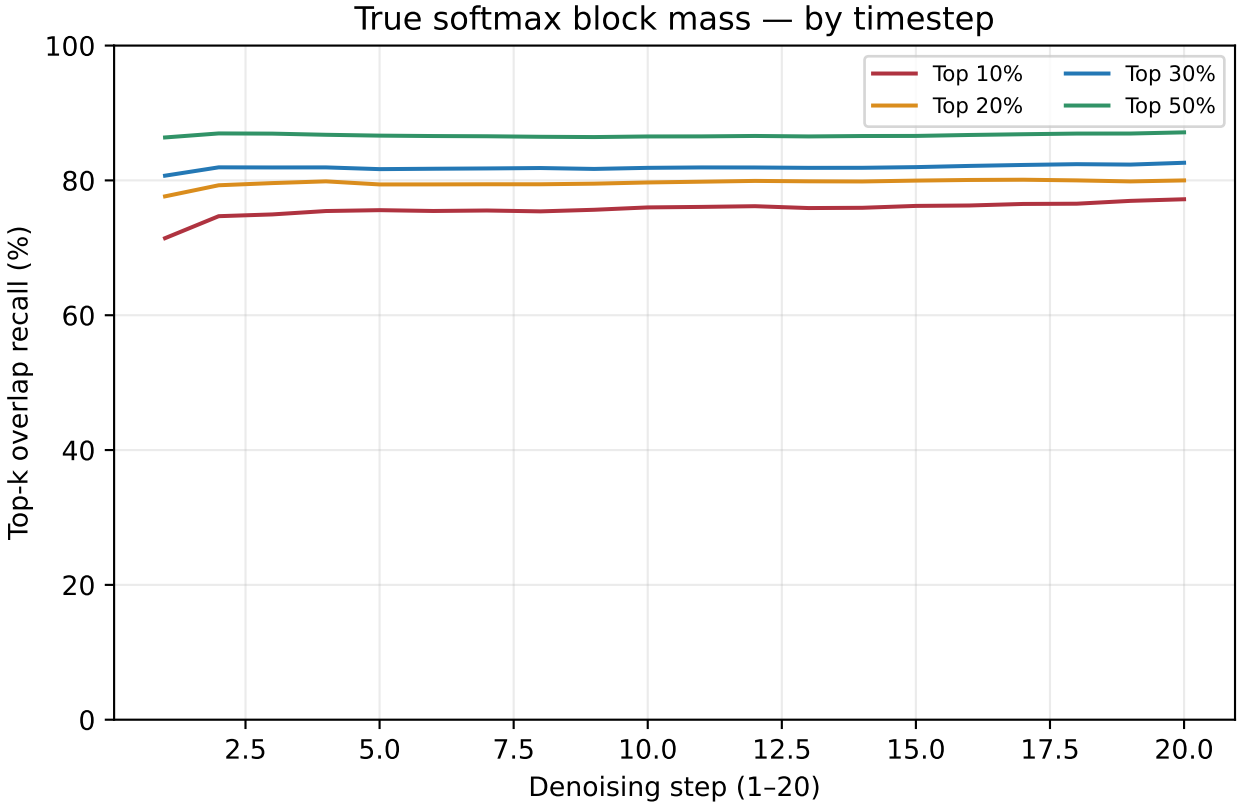

True softmax block mass — by layer depth

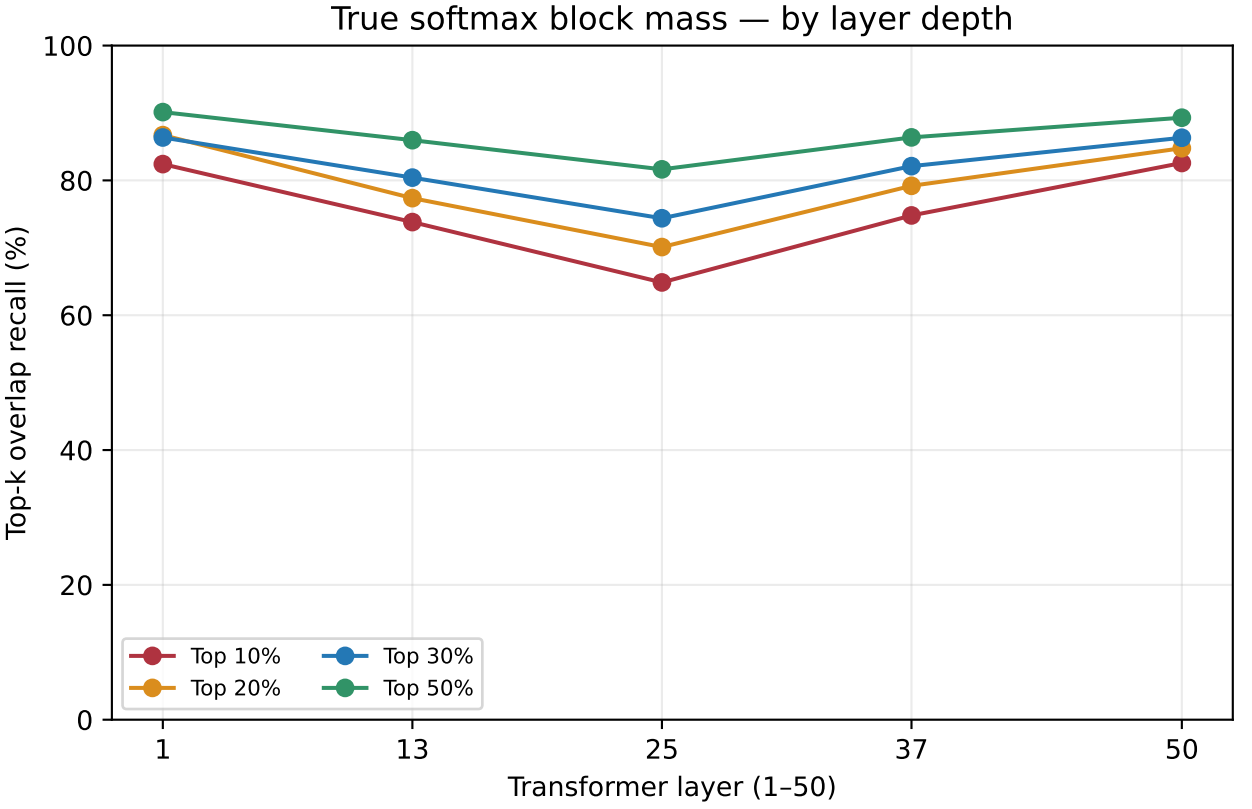

Block max QK logit — by timestep

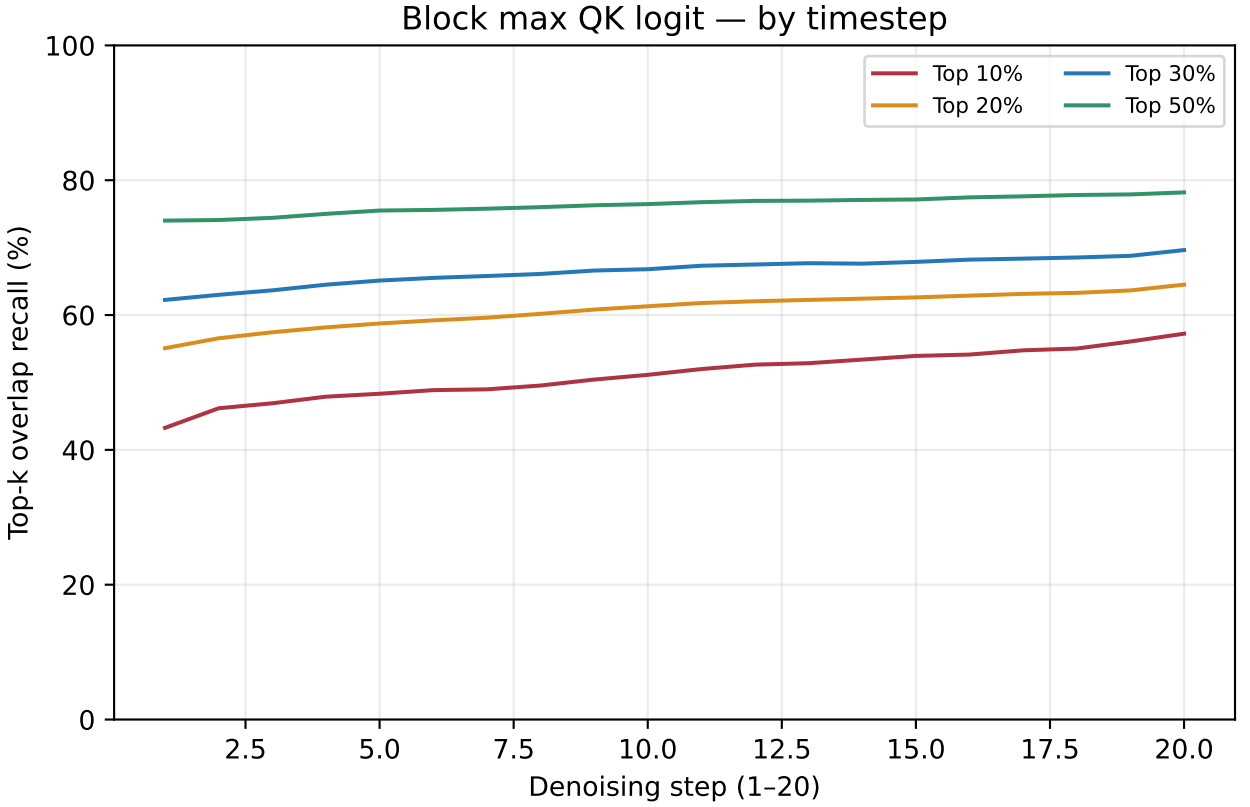

Block max QK logit — by layer depth

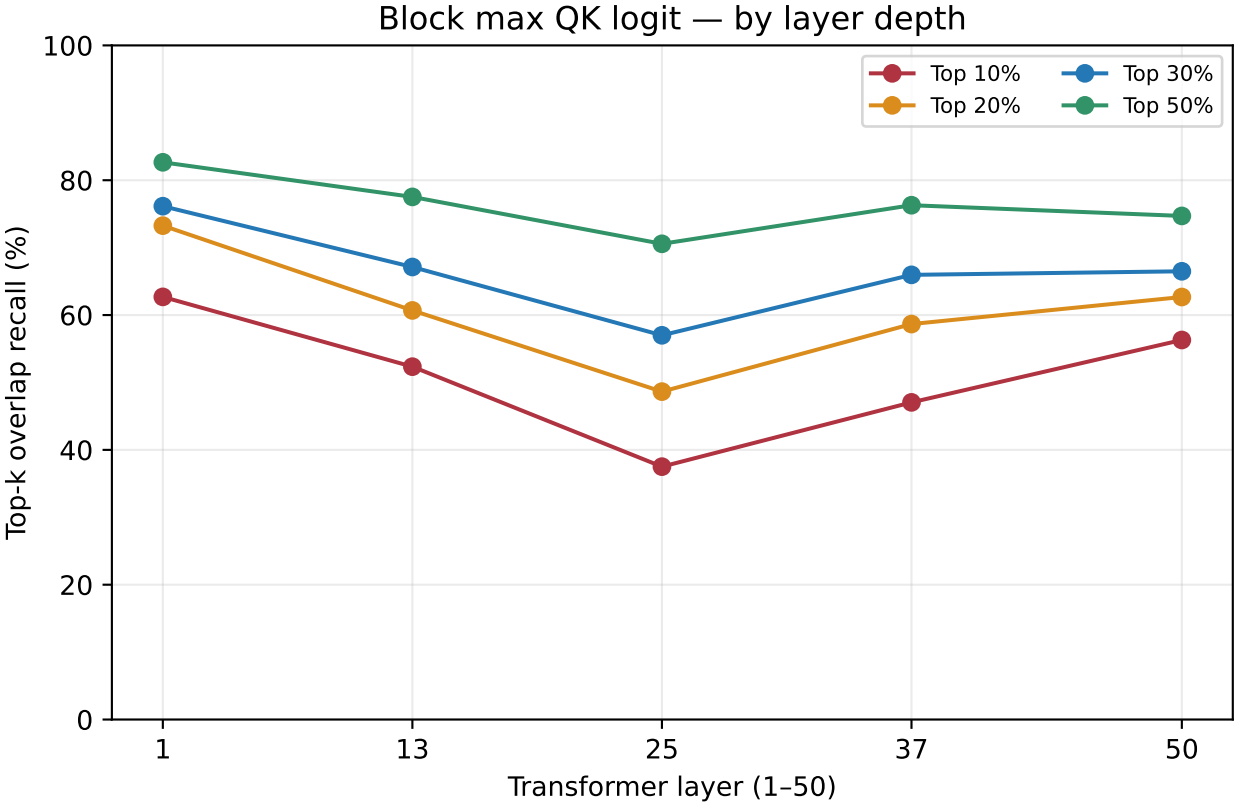

Block max softmax probability — by timestep

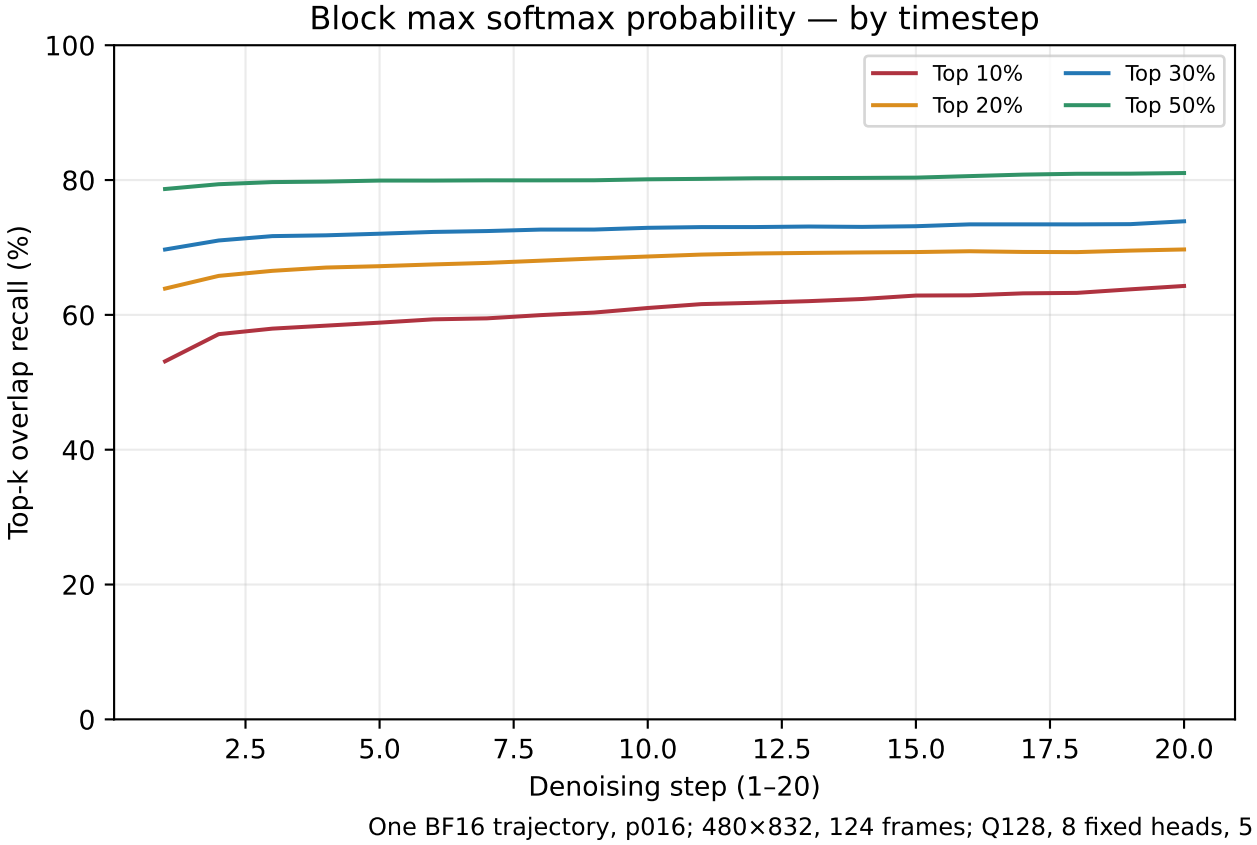

Block max softmax probability — by layer depth

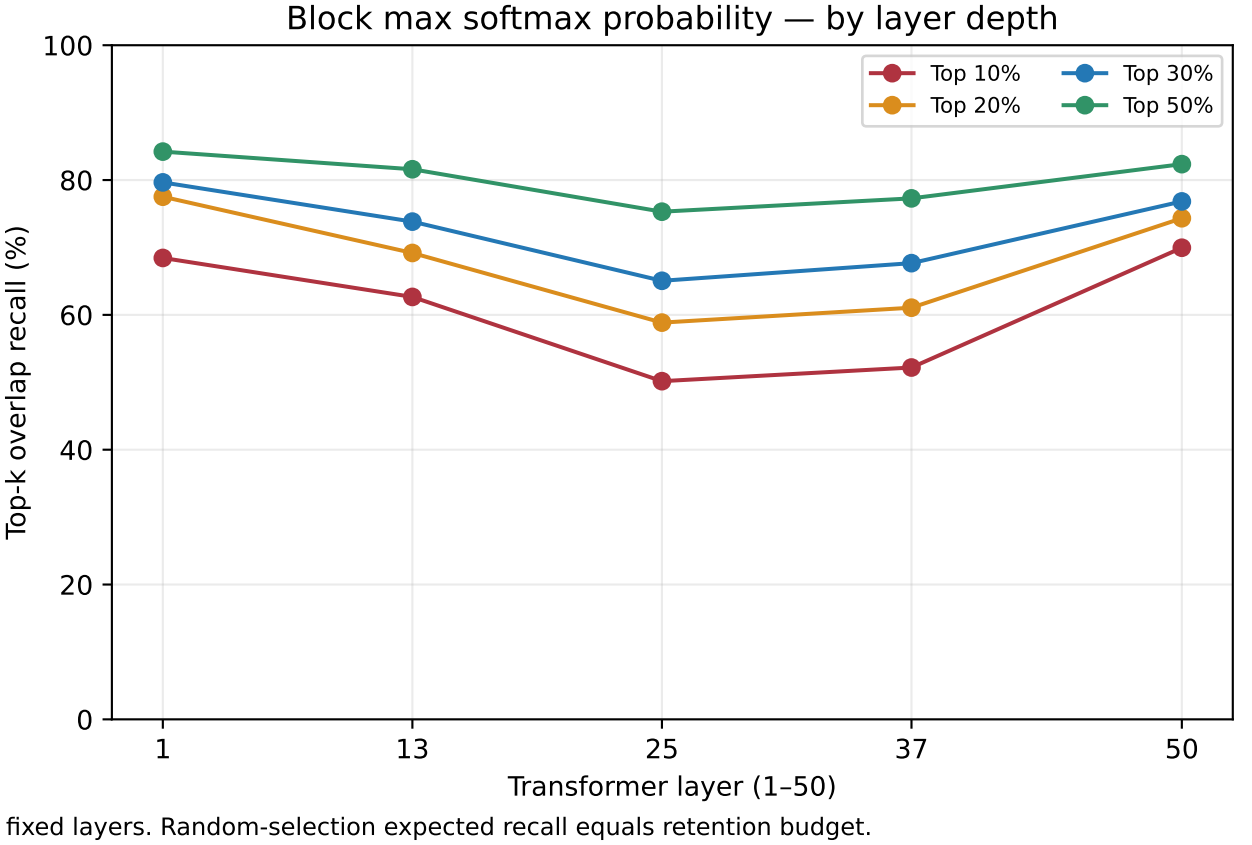

One BF16 trajectory, p016; 480×832, 124 frames; Q128, 8 fixed heads, 5 fixed layers. Random-selection expected recall equals retention budget.